Homework5 code

November 24, 2024

1 Spam ham classifier

Consider a text classification problem. In this case, you will try to classify text as either spam or ham. To do this, you will apply concepts of Likelihood, prior, and posterior given a dataset comprising pairs of text and labels. There are two types of labels: 1 (spam) and 0 (ham). Your goal is to create a simple classifier that, when given, determines if the text is spam or ham.

[1]:
```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import nltk
from nltk.corpus import stopwords
from sklearn.metrics import accuracy_score
import re
```

1.1 Data loading and cleaning

[2]:
```
data = pd.read_csv('spam_ham_dataset.csv')
data = data.dropna()
data = data.drop_duplicates()
df =data
nltk.download('stopwords')
stop = stopwords.words('english')
```

```
[nltk_data] Downloading package stopwords to /home/kip/nltk_data…
[nltk_data]   Package stopwords is already up-to-date!
```

[ ]:
```
def remove_stopwords(text):
    text = [word.lower() for word in text.split() if word.lower() not in stop]
    # remove special characters
    text = [re.sub('\W+','', word) for word in text]

    return " ".join(text)

txt1 = df.text[0]
```

```python
print(txt1)
# remove the filler words
txt1 = remove_stopwords(txt1)
print(txt1)
# apply the function to the entire dataset
df['text'] = df['text'].apply(remove_stopwords)
df.head()
```

```
<>:4: SyntaxWarning: invalid escape sequence '\W'
<>:4: SyntaxWarning: invalid escape sequence '\W'
/tmp/ipykernel_1560/2336932942.py:4: SyntaxWarning: invalid escape sequence '\W'
  text = [re.sub('\W+','', word) for word in text]
```

```
Subject: enron methanol ; meter # : 988291
this is a follow up to the note i gave you on monday , 4 / 3 / 00 { preliminary
flow data provided by daren } .
please override pop ' s daily volume { presently zero } to reflect daily
activity you can obtain from gas control .
this change is needed asap for economics purposes .
subject enron methanol  meter   988291 follow note gave monday  4  3  00
preliminary flow data provided daren   please override pop  daily volume
presently zero  reflect daily activity obtain gas control  change needed asap
economics purposes
```

[ ]:

```
   Unnamed: 0 label                                               text  \
0         605   ham  subject enron methanol  meter   988291 follow ...
1        2349   ham  subject hpl nom january 9  2001  see attached ...
2        3624   ham  subject neon retreat ho ho ho   around wonderf...
3        4685  spam  subject photoshop windows  office  cheap  mai...
4        2030   ham  subject  indian springs deal book teco pvr rev...

   label_num
0          0
1          0
2          0
3          1
4          0
```

Next we split the data into training and testing. We will derive the probabilities from the training data and then use them to predict the testing data.

[3]:

```python
# lets split the data into training and testing
X = df.text
y = df.label_num
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
 ↪random_state=42)
```

1.1.1 (1)

Find the priors. What are the priors in this distribution? i.e find P (ham) and P (spam)

[4]:
```
# Calculate priors
P_ham = y_train.value_counts()[0] / len(y_train)
P_spam = y_train.value_counts()[1] / len(y_train)

print(f"P(ham): {P_ham}")
print(f"P(spam): {P_spam}")
```

```
P(ham): 0.7084139264990329
P(spam): 0.2915860735009671
```

1.1.2 (2)

Find the likelihoods for each word. For each word in the dataset, find the likelihood that the word is in spam and ham. This will represent the conditional probability $P(w|spam)$ and $P(w|ham)$ for w where $w \quad V$. V is the vocabulary of the dataset.

[5]:
```
from collections import defaultdict, Counter

# Initialize counters for spam and ham words
spam_words = Counter()
ham_words = Counter()

# Separate spam and ham texts
spam_texts = X_train[y_train == 1]
ham_texts = X_train[y_train == 0]

# Count words in spam and ham texts
for text in spam_texts:
    for word in text.split():
        spam_words[word] += 1

for text in ham_texts:
    for word in text.split():
        ham_words[word] += 1

# Calculate total number of words in spam and ham texts
total_spam_words = sum(spam_words.values())
total_ham_words = sum(ham_words.values())

# Calculate likelihoods
likelihoods_spam = {word: (count / total_spam_words) for word, count in␣
 ↪spam_words.items()}
likelihoods_ham = {word: (count / total_ham_words) for word, count in ham_words.
 ↪items()}
```

```
print("Likelihoods for spam words:", list(likelihoods_spam.items())[:10])
print("Likelihoods for ham words:", list(likelihoods_ham.items())[:10])
```

```
Likelihoods for spam words: [('Subject:', 0.004289662482526562), ('message',
0.0008536641756271764), ('subject', 0.0004446167581391544), ('hey',
0.00011737882414873675), ('i', 0.003272379339904176), ("'",
0.005082858779046813), ('am', 0.00035213647244621026), ('julie',
7.11386813022647e-06), ('^', 0.0001316065604091897), ('_',
0.0026997129554209454)]
Likelihoods for ham words: [('Subject:', 0.0044757218820028625), ('april',
0.0006201853529328199), ('activity', 0.00032078552737904476), ('surveys',
1.9858151694893248e-05), ('we', 0.0038372059275055257), ('are',
0.0026579372268549424), ('starting', 7.637750651882018e-05), ('to',
0.01828935771099668), ('collect', 1.8330601564516844e-05), ('data',
0.0002520457715121066)]
```

1.1.3 (3)

Define a function that, when given a text sequence, returns the probability of the text being in spam. I.e., it returns P (spam|text). Note that this function calculates the likelihood using the Bayes rule. Do the same for ham.

```
[6]: def calculate_posterior(text, priors, likelihoods, total_words):
    '''
    Calculate the posterior probability of a text being spam or ham
    given the text and the likelihoods of spam and ham words.

    Parameters:
    text (str): the text to classify
    priors (tuple): the prior probabilities of spam and ham
    likelihoods (dict): the likelihoods of spam and ham words
    total_words (int): the total number of words in the training set

    returns:
    float: the posterior probability of the text being spam or ham
    '''

    # Split the text into words
    words = text.split()
    # Initialize posterior as the log of the priors
    posterior = np.log(priors)

    # Calculate the posterior for spam and ham of the text
    for word in words:
        # If the word is in the likelihoods dictionary, add the log likelihood␣
↪to the posterior
        if word in likelihoods:
            posterior += np.log(likelihoods[word])
```

```python
        #  If the word is not in the likelihoods dictionary, apply Laplace␣
↪smoothing to avoid zero probabilities
        else:
            # Apply Laplace smoothing for unseen words
            posterior += np.log(1 / (total_words + len(likelihoods)))

    # Return the final posterior probability
    return posterior

def predict_spam(text):
    '''
    Predict whether a text is spam given the text.

    Parameters:
    text (str): the text to classify

    returns:
    float: the posterior probability of the text being spam
    '''
    P_spam_given_text = calculate_posterior(text, P_spam, likelihoods_spam,␣
↪total_spam_words)
    return P_spam_given_text

def predict_ham(text):
    '''
    Predict whether a text is ham given the text.

    Parameters:
    text (str): the text to classify

    returns:
    float: the posterior probability of the text being ham
    '''
    P_ham_given_text = calculate_posterior(text, P_ham, likelihoods_ham,␣
↪total_ham_words)
    return P_ham_given_text

# Example usage
text_example = "Congratulations, spin and win money now"
print(f"P(spam|text): {predict_spam(text_example)}")
print(f"P(ham|text): {predict_ham(text_example)}")
```

```
P(spam|text): -54.51657329606763
P(ham|text): -60.754027884439054
```

1.1.4 (4)

Perform inference, i.e., given a string of text, determine if it is ham or spam based on the posterior probabilities calculated from the previous steps. Your function will determine the posterior probability of your text being in ham and spam and classify it as being the larger of the two.

[8]:
```
def classify_text(text):
    P_spam_given_text = predict_spam(text)
    P_ham_given_text = predict_ham(text)

    if P_spam_given_text > P_ham_given_text:
        return 'spam'
    else:
        return 'ham'

# Example usage
text_example = "Congratulations, spin and win money now"
classification = classify_text(text_example)
print(f"The text '{text_example}' is classified as: {classification}")
```

```
The text 'Congratulations, spin and win money now' is classified as: spam
```

1.1.5 (5)

Evaluate the data based on your test set and report the accuracy of your classifier. Your accuracy must be greater than 85%.

[11]:
```
# Predict the labels for the test set
y_pred = X_test.apply(classify_text)

# Convert predictions to numerical labels
y_pred_num = y_pred.apply(lambda x: 1 if x == 'spam' else 0)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred_num)
print(f"Accuracy: {accuracy * 100:.2f}%")
```

```
Accuracy: 96.52%
```