Homework 1 - Introduction to Probabilistic Graphical Models

kipngeno koech - bkoech

February 10, 2025

1 Bayesian Networks

1. Consider a simple Markov Chain structure $X \to Y \to Z$, where all variables are binary. You are required to:

   (a) Write a code (using your preferred programming language) that generates a distribution (not necessarily a valid BN one) over the 3 variables.

   [**in the notebook**]

   (b) Write a code that verifies whether a distribution is a valid BN distribution.

   [**in the notebook**]

   (c) Using your code, generate 10000 distributions and compute the fraction of distributions that are valid BN distributions.

   [**in the notebook**]
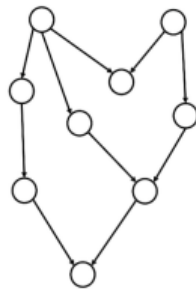
2. Given the following Bayesian Network

Figure 1: A Bayesian network.

Figure 1: Bayesian Network

   (a) Propose a topological ordering of this graph
       In Figure 2, the topological ordering is:
       i. $A \to B \to C \to D \to E \to F \to G \to H \to I$
       ii. $A \to B \to C \to E \to D \to F \to G \to I \to H$

   (b) Let $\mathbf{X}$ be a random vector that is Markov with respect to the graph. We assume that the random variables X, are binary. Write all the local conditional independence

$$X_A \text{ has no parents, so no independence condition applies here.}$$

$$X_B \text{ has no parents, so no independence condition applies here.}$$

X_C is conditionally independent of all other nodes given its parent X_A :

$$X_C \perp \{X_B, X_D, X_E, X_F, X_G, X_H, X_I\} \mid X_A$$

X_D is conditionally independent of all other nodes given its parent X_A :

$$X_D \perp \{X_B, X_C, X_E, X_F, X_G, X_H, X_I\} \mid X_A$$

X_E is conditionally independent of all other nodes given its parents X_C and X_D :

$$X_E \perp \{X_A, X_B, X_F, X_G, X_H, X_I\} \mid \{X_C, X_D\}$$

X_F is conditionally independent of all other nodes given its parent X_B :

$$X_F \perp \{X_A, X_C, X_D, X_E, X_G, X_H, X_I\} \mid X_B$$

X_G is conditionally independent of all other nodes given its parents X_E and X_F :

$$X_G \perp \{X_A, X_B, X_C, X_D, X_H, X_I\} \mid \{X_E, X_F\}$$

X_H is conditionally independent of all other nodes given its parent X_D :

$$X_H \perp \{X_A, X_B, X_C, X_E, X_F, X_G, X_I\} \mid X_D$$

X_I is conditionally independent of all other nodes given its parents X_G and X_H :

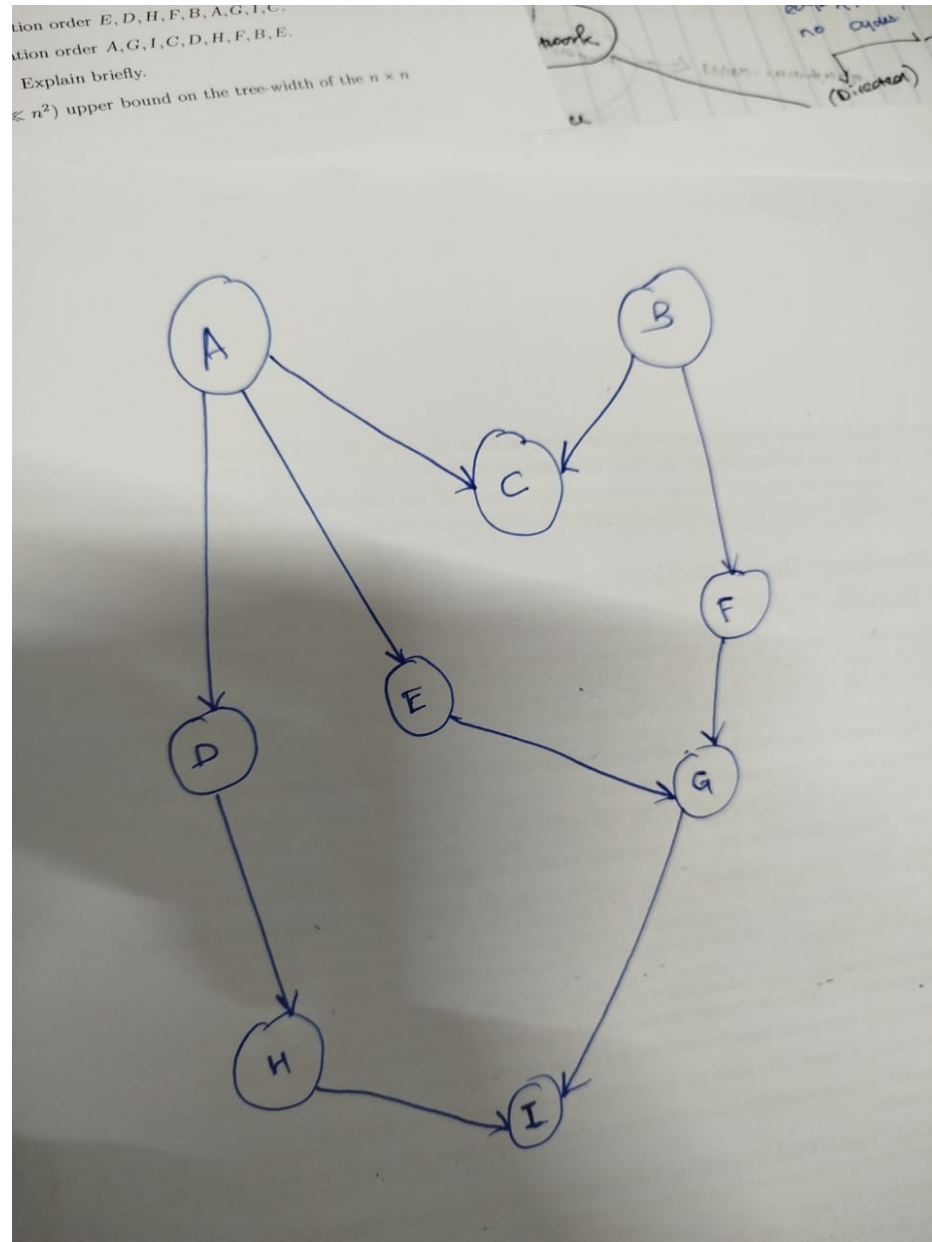$$X_I \perp \{X_A, X_B, X_C, X_D, X_E, X_F\} \mid \{X_G, X_H\}$$

Figure 2: Bayesian Network

3. State True or False, and briefly justify your answer within 3 lines. The statements are either direct consequences of theorems in Koller and Friedman (2009, Ch. 3), or have a short proof. In the follows, P is a distribution and G is a BN structure.

(a) If $A \perp B \mid C$ and $A \perp C \mid B$, then $A \perp B$ and $A \perp C$. (Suppose the joint distribution of A, B, C is positive.) (This is a general probability question not related to BNs.)

- **False**. Whilst local independence implies global independence, global independence does not imply local independence.
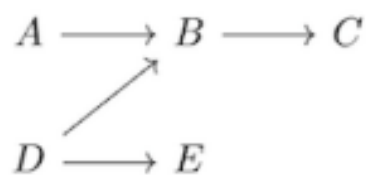

Figure 2: A Bayesian network.

Figure 3: Bayesian Network

(b) In Figure 2, $E \perp C \mid B$
   True. E is conditionally independent of C given B because we have conditioned on the node B.

(c) in Figure 2, $A \perp E \mid C$
   False. A is not conditionally independent of E given C because there is a path from A to E that is not blocked by C.

In figure 3, Recall the definitions of local and global independences of G and independences of P.

$$I_l(G) = \{(X \perp \text{NonDescendants}_G(X) \mid \text{Parents}_G(X))\} \tag{1}$$

$$I(G) = \{(X \perp Y \mid Z) : \text{d-separated}_G(X, Y \mid Z)\} \tag{2}$$

$$I(P) = \{(X \perp Y \mid Z) : P(X, Y \mid Z) = P(X \mid Z)P(Y \mid Z)\} \tag{3}$$

(d) In Figure 3, relation 1 is true.

(e) In Figure 3, relation 2 is true.

(f) In Figure 3, relation 3 is true.

(g) If G is an I-map for P, then P may have extra conditional independencies than G.

$$P \text{ factorizes over } \mathcal{G} \xrightarrow{\ (1)\ } \mathcal{I}(\mathcal{G}) \subseteq \mathcal{I}(P) \xrightarrow{\ (2)\ } \mathcal{I}_\ell(\mathcal{G}) \subseteq \mathcal{I}(P)$$
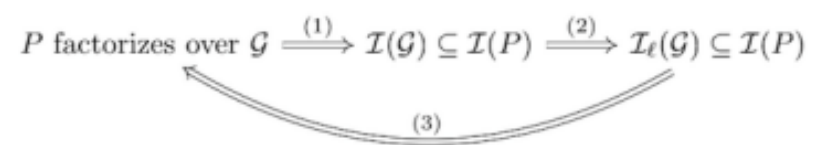
<p align="center">(3)</p>

Figure 3: Some relations in Bayesian networks.

(h) Two BN structures G_1 and G_2 are I-equivalent if they have the same skeleton and the same set of v-structures.

(i) If G_1 is an I-map of distribution P, and G_1 has fewer edges than G_2, then G_2 is not a minimal I-map of P.

(j) The P-map of a distribution, if it exists, is unique.

2 Markov Networks

Let $\mathbf{X} = (X_1, \dots, X_d)$ be a random vector with mean $\boldsymbol{\mu}$ and covariance matrix $\boldsymbol{\Sigma}$. The partial correlation matrix $\mathbf{R}$ of $\mathbf{X}$ is a $d \times d$ matrix where each entry $R_{ij} = \rho(X_i, X_j \mid \mathbf{X}_{-ij})$ is the partial correlation between X_i and X_j given the $d-2$ remaining variables $\mathbf{X}_{-ij}$. Let $\boldsymbol{\Theta} = \boldsymbol{\Sigma}^{-1}$ be the inverse covariance matrix of $\mathbf{X}$.

We will prove the relation between $\mathbf{R}$ and $\boldsymbol{\Theta}$, and furthermore how $\boldsymbol{\Theta}$ characterizes conditional independence in Gaussian graphical models.

1. (10 points) Show that

$$\begin{pmatrix} \Theta_{ii} & \Theta_{ij} \\ \Theta_{ji} & \Theta_{jj} \end{pmatrix} = \begin{pmatrix} \mathrm{Var}[e_i] & \mathrm{Cov}[e_i, e_j] \\ \mathrm{Cov}[e_i, e_j] & \mathrm{Var}[e_j] \end{pmatrix}^{-1}$$

for any $i, j \in [d], i \neq j$. Here e_i is the residual resulting from the linear regression of X_{-ij} to X_i, and similarly e_j is the residual resulting from the linear regression of X_{-ij} to X_j.

The residuals e_i and e_j are given by:

$$e_i = X_i - \mathbb{E}[X_i | X_{-ij}], \quad e_j = X_j - \mathbb{E}[X_j | X_{-ij}].$$

These residuals are uncorrelated with X_{-ij}, meaning $\mathrm{Cov}(e_i, X_{-ij}) = 0$ and $\mathrm{Cov}(e_j, X_{-ij}) = 0$.

Covariance of Residuals: The covariance matrix of the residuals e_i and e_j is:

$$\mathrm{Cov}(e_i, e_j) = \mathrm{Cov}(X_i, X_j | X_{-ij}).$$

This is because the residuals capture the conditional covariance between X_i and X_j given X_{-ij}.

Inverse Covariance Matrix: The inverse covariance matrix $\Theta = \Sigma^{-1}$ satisfies:

$$\Theta = \begin{pmatrix} \Theta_{ii} & \Theta_{ij} \\ \Theta_{ji} & \Theta_{jj} \end{pmatrix}.$$

By the properties of the inverse covariance matrix, the conditional covariance matrix of (X_i, X_j) given X_{-ij} is:

$$\mathrm{Cov}(X_i, X_j | X_{-ij}) = \begin{pmatrix} \Theta_{ii} & \Theta_{ij} \\ \Theta_{ji} & \Theta_{jj} \end{pmatrix}^{-1}.$$

Equating the Matrices: Since $\mathrm{Cov}(e_i, e_j) = \mathrm{Cov}(X_i, X_j | X_{-ij})$, we have:

$$\begin{pmatrix} \Theta_{ii} & \Theta_{ij} \\ \Theta_{ji} & \Theta_{jj} \end{pmatrix} = \begin{pmatrix} \mathrm{Var}[e_i] & \mathrm{Cov}[e_i, e_j] \\ \mathrm{Cov}[e_i, e_j] & \mathrm{Var}[e_j] \end{pmatrix}^{-1}.$$

2. (10 points) Show that

$$R_{ij} = -\frac{\Theta_{ij}}{\sqrt{\Theta_{ii}\Theta_{jj}}}$$

Partial Correlation: The partial correlation R_{ij} is defined as:

$$R_{ij} = \rho(X_i, X_j | X_{-ij}) = \frac{\mathrm{Cov}(X_i, X_j | X_{-ij})}{\sqrt{\mathrm{Var}(X_i | X_{-ij})\mathrm{Var}(X_j | X_{-ij})}}.$$

Conditional Covariance and Variance: From Problem 1, we know:

$$\mathrm{Cov}(X_i, X_j | X_{-ij}) = -\frac{\Theta_{ij}}{\Theta_{ii}\Theta_{jj} - \Theta_{ij}^2},$$

and:

$$\mathrm{Var}(X_i | X_{-ij}) = \frac{1}{\Theta_{ii}}, \quad \mathrm{Var}(X_j | X_{-ij}) = \frac{1}{\Theta_{jj}}.$$

Substitute into Partial Correlation: Substituting these into the definition of R_{ij}, we get:

$$R_{ij} = \frac{-\frac{\Theta_{ij}}{\Theta_{ii}\Theta_{jj} - \Theta_{ij}^2}}{\sqrt{\frac{1}{\Theta_{ii}} \cdot \frac{1}{\Theta_{jj}}}}.$$

Simplifying, we obtain:

$$R_{ij} = -\frac{\Theta_{ij}}{\sqrt{\Theta_{ii}\Theta_{jj}}}.$$

3. (15 points) From the above result and the relation between independence and correlation, we know

$$\Theta_{ij} = 0 \iff R_{ij} = 0 \implies X_i \perp X_j \mid X_{-ij}$$

Note the last implication only holds in one direction. Now suppose $X \sim N(\mu, \Sigma)$ is jointly Gaussian. Show that $R_{ij} = 0 \implies X_i \perp X_j \mid X_{-ij}$.

Partial Correlation and Conditional Independence: From Problem 2, we know:

$$R_{ij} = -\frac{\Theta_{ij}}{\sqrt{\Theta_{ii}\Theta_{jj}}}.$$

If $R_{ij} = 0$, then $\Theta_{ij} = 0$.

Inverse Covariance and Conditional Independence: For a jointly Gaussian distribution, the inverse covariance matrix $\Theta = \Sigma^{-1}$ encodes conditional independence. Specifically:

$$\Theta_{ij} = 0 \iff X_i \perp X_j | X_{-ij}.$$

This is because the off-diagonal elements of Θ represent the conditional dependence between variables after accounting for all other variables.

Conclusion: Since $R_{ij} = 0$ implies $\Theta_{ij} = 0$, it follows that:

$$X_i \perp X_j | X_{-ij}.$$

3 Exact Inference - Variable Elimination

Reference materials for this problem:

- Jordan textbook Ch. 3, available at `https://people.eecs.berkeley.edu/~jordan/prelims/chapter3.pdf`

- Koller and Friedman (2009, Ch. 9 and Ch. 10)

3.1 Variable elimination on a grid [10 points]

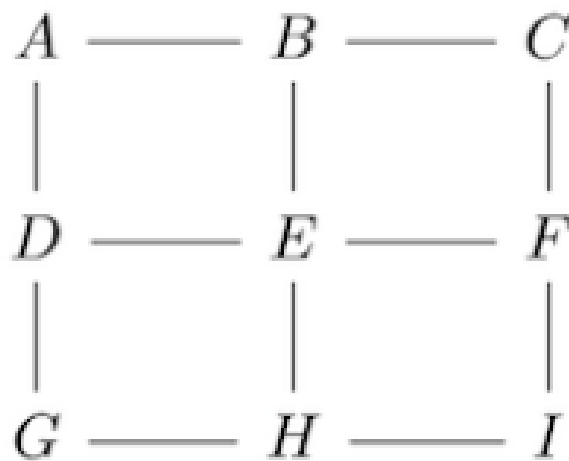Consider the following Markov network:

Figure 4: Markov Network

We are going to see how tree-width, a property of the graph, is related to the intrinsic complexity of variable elimination of a distribution

1. (5 points) Write down largest clique(s) for the elimination order $E, D, H, F, B, A, G, I, C$.
   We start by eliminating E, its neighbors are D, F, H, B, so the clique here is :$\{D, F, H, B\}$. size: **4**
   We then eliminate D, its neighbors are A, B, F, G, H, so the clique here is: $\{A, B, f, G, H\}$. size: **5**
   We then eliminate H, its neighbors are G, A, I, B, F, so the clique here is: $\{G, A, I, B, F\}$. size: **5**
   We then eliminate F, its neighbors are A, B, c, G, I, so the clique here is: $\{A, B, C, G, I\}$. size: **5**
   We then eliminate B, its neighbors are A, C, G, I, so the clique here is: $\{A, C, G, I\}$. size: **4**
   We then eliminate A, its neighbors are C, G, I, so the clique here is: $\{C, G, I\}$. size: **3**
   We then eliminate G, its neighbors are C, I, so the clique here is: $\{C, I\}$. size: **2**
   We then eliminate I, its neighbors are C, so the clique here is: $\{C\}$. size: **1**
   The largest clique(s) for the elimination order $E, D, H, F, B, A, G, I, C$ is **5**

2. (5 points) Write down largest clique(s) for the elimination order $A, G, I, C, D, H, F, B, E$.
   We start by eliminating A, its neighbors are D, B, so the clique here is: $\{D, B\}$. size: **2**
   We then eliminate G, its neighbors are D, H, so the clique here is: $\{D, H\}$. size: **2**
   We then eliminate I, its neighbors are H, F, so the clique here is: $\{H, F\}$. size: **2**
   We then eliminate C, its neighbors are F, B, so the clique here is: $\{F, B\}$. size: **2**
   We then eliminate D, its neighbors are H, B, E, so the clique here is: $\{H, B, E\}$. size: **3**
   We then eliminate H, its neighbors are F, B, E, so the clique here is: $\{F, B, E\}$. size: **3**
   We then eliminate F, its neighbors are B, E, so the clique here is: $\{B, E\}$. size: **2**
   We then eliminate B, its neighbors are E, so the clique here is: $\{E\}$. size: **1**
   The largest clique(s) for the elimination order $A, G, I, C, D, H, F, B, E$ is **3**

3. (5 points) Which of the above ordering is preferable? Explain briefly.
   The second ordering is preferable because it results in a smaller largest clique size (3) compared to the first ordering (5). A smaller clique size generally leads to more efficient computations in variable elimination algorithms.

4. (10 points) Using this intuition, give a reasonable ($\ll n^2$) upper bound on the tree-width of the $n \times n$ grid.
   The tree-width of an $n \times n$ grid is at most $2n - 2$. This is because the grid can be decomposed into a tree structure where each node has at most $2n - 2$ neighbors. This upper bound is reasonable and much smaller than n^2, which would be the case for a fully connected graph.