

Homework 3 - Introduction to Probabilistic Graphical Models

kipngeno koech - bkoech

1 Sampling (Monte-Carlo), MCMC[50 points]

1.1 (Monte-Carlo) [20 points]

Given a random distribution $p(x)$ on $x = [x_1, \dots, x_D]^T \in \mathbb{R}^D$. Suppose we want to perform inference $\mathbb{E}_{p(x)}[f(x)]$ using importance sampling, with $q(x)$ as the proposal distribution. According to importance sampling, we draw L i.i.d. samples $x^{(1)}, \dots, x^{(L)}$ from $q(x)$, and we have

$$\mathbb{E}_{p(x)}[f(x)] \approx \frac{1}{\sum_{i=1}^L u_i} \sum_{i=1}^L f(x^{(i)}) u_i$$

where the (unnormalized) importance weights $u_i = \frac{p(x^{(i)})}{q(x^{(i)})}$.

1. (5 points) Find the mean and variance of the unnormalized importance weights $\mathbb{E}_{q(x)}[u_i]$ and $\text{Var}_{q(x)}[u_i]$.
2. (5 points) Prove the following lemma: $\mathbb{E}_{p(x)} \left[\frac{p(x)}{q(x)} \right] \geq 1$, and the equality holds only when $p = q$.
3. (9 points) A measure of the variability of two components in vector $u = [u_1, \dots, u_L]^T$ is given by $\mathbb{E}_{q(x)} [(u_i - u_j)^2]$. Assume that both p and q can be factorized, i.e. $p(x) = \prod_{i=1}^D p_i(x_i)$, and $q(x) = \prod_{i=1}^D q_i(x_i)$. Show that $\mathbb{E}_{q(x)} [(u_i - u_j)^2]$ has exponential growth with respect to D .
4. (1 point) Use the conclusion in (c) to explain why the standard importance sampling does not scale well with dimensionality and would blow up in high-dimensional cases.

1.2 (MCMC)[40 points]

1.2.1 Multiple Choice [10 points]

1. (5 points) Which of the following statements is true for the acceptance probability $A(x'|x) = \min(1, \frac{P(x')Q(x|x')}{P(x)Q(x'|x)})$ of Metropolis-Hastings algorithm?

- (a) It satisfies detailed balance.
 - (b) We can just evaluate $P(x')$ and $P(x)$ up to a normalization constant.
 - (c) It ensures that the MH algorithm eventually converges to the true distribution.
 - (d) All of the above.
2. (5 points) Which of the following statements is true for Hamiltonian Monte Carlo in comparison with vanilla MCMC?
- (a) It can improve acceptance rate and give better mixing.
 - (b) Stochastic variants can be used to improve performance in large dataset scenarios.
 - (c) It may not be used for discrete variable.
 - (d) All of the above.

1.2.2 Modeling with Markov Chain Monte Carlo [30 points]

We are going to use the data from the 2013-2014 Premier League PL1 to build a predictive model on the number of goals scored in a single game by the two opponents. Bayesian hierarchical model is a good candidate for this kind of modeling task. We model each team's strength (both attacking and defending) as latent variables. Then in each game, the goals scored by the home team is a random variable conditioned on the attacking strength of the home team and the defending strength of the away team. Similarly, the goals scored by the away team is a random variable conditioned on the attack strength of the away team and the defense strength of the home team. Therefore, the distribution of the scoreline of a specific game is dependent on the relative strength between the home team A and the away team B, which also depends on the relative strength between those teams with their other opponents.

Table 1: 2013-2014 Premier League teams

Index	0	1	2	3
Team	Arsenal	Aston Villa	Cardiff City	Chelsea
Index	4	5	6	7
Team	Crystal Palace	Everton	Fulham	Hull City
Index	8	9	10	11
Team	Liverpool	Manchester City	Manchester United	Newcastle United
Index	12	13	14	15
Team	Norwich City	Southampton	Stoke City	Sunderland
Index	16	17	18	19
Team	Swansea City	Tottenham Hotspur	West Bromwich Albion	West Ham United

Here we consider using the same model as described by Baio and Blangiardo. The Premier League has 20 teams, and we index them as in Table 1. Each team

would play 38 matches every season (playing each of the other 19 teams home and away), which totals 380 games in the entire season. For the g -th game, assume that the index of home team is $h(g)$ and the index of the away team is $a(g)$. the observed number of goals is:

$$y_{gj}|\theta_{gj} = \text{Poisson}(\theta_{gj})$$

where the $\theta = (\theta_{g1}, \theta_{g2})$ represent the scoring intensity in the g -th game for the team playing at home ($j = 1$) and away ($j = 2$), respectively. We put a log-linear model for the θ s:

$$\log \theta_{g1} = \text{home} + \text{att}_{h(g)} - \text{def}_{a(g)}$$

$$\log \theta_{g2} = \text{att}_{a(g)} - \text{def}_{h(g)}$$

Note that team strength is broken into attacking and defending strength. And home represents home- team advantage, and in this model is assumed to be constant across teams. The prior on the home is a normal distribution

$$\text{home} \sim \mathcal{N}(0, \tau_0^{-1})$$

where the precision $\tau_0 = 0.0001$.

The team-specific attacking and defending effects are modeled as exchangeable:

$$\text{att}_t \sim \mathcal{N}(\mu_{\text{att}}, \tau_{\text{att}}^{-1})$$

$$\text{def}_t \sim \mathcal{N}(\mu_{\text{def}}, \tau_{\text{def}}^{-1})$$

We use conjugate priors as the hyper-priors on the attack and defense means and precisions:

$$\mu_{\text{att}} \sim \mathcal{N}(0, \tau_1^{-1})$$

$$\mu_{\text{def}} \sim \mathcal{N}(0, \tau_1^{-1})$$

$$\tau_{\text{att}} \sim \text{Gamma}(\alpha, \beta)$$

$$\tau_{\text{def}} \sim \text{Gamma}(\alpha, \beta)$$

where the precision $\tau_1 = 0.0001$, and we set parameters $\alpha = \beta = 0.1$.

This hierarchical Bayesian model can be represented using a directed acyclic graph as shown in Figure 1. Where the goals of each game $y = \{y_{gj}|g = 0, \dots, 379, j = 1, 2\}$ are 760 observed variables, and parameters $\theta = (\text{home}, \text{att}_0, \dots, \text{att}_{19}, \text{def}_0, \dots, \text{def}_{19})$ and hyper-parameters $\eta = (\mu_{\text{att}}, \mu_{\text{def}}, \tau_{\text{att}}, \tau_{\text{def}})$ are unobserved variables that we need to make inference. To ensure identifiability, we enforce a corner constraint on the parameters (pinning one team's parameters to 0,0). Here we use the first team as reference and assign its attacking and defending strength to be 0:

$$\text{att}_0 = \text{def}_0 = 0$$

In this question, we want to estimate the posterior mean of the attacking and defending strength for each team, i.e. $\mathbb{E}_{p(\theta, \eta|y)}[\text{att}_i]$, $\mathbb{E}_{p(\theta, \eta|y)}[\text{def}_i]$, and $\mathbb{E}_{p(\theta, \eta|y)}[\text{home}]$.

1. (10 points) Find the joint likelihood $p(y, \theta, \eta)$.
2. (10 points) Write down the Metropolis-Hastings algorithm for sampling from posterior $p(\theta, \eta|y)$, and derive the acceptance function for a proposal distribution of your choice (e.g. isotropic Gaussian).
3. (10 points) Implement the M-H algorithm to inference the posterior distribution. The data can be found from `premier_league_2013.2014.dat`, which contains a 380×4 matrix. The first column is the number of goals y_{g1} scored by the home team, the second column is the number of goals y_{g2} scored by the away team, the third column is the index for the home team $h(g)$, and the fourth column is the index for the away team $a(g)$. Use isotropic Gaussian proposal distribution, $\mathcal{N}(0, \sigma^2 I)$ and 0 as the starting point. Run the MCMC chain for 5000 steps to burn in and then collect 5000 samples with t steps in between (i.e., run M-H for $5000t$ steps and collect only each t -th sample). This is called thinning, which reduces the autocorrelation of the MCMC samples introduced by the Markovian process. The parameter sets are $\sigma = 0.05$, and $t = 5$. Plot the trace plot of the burning phase and the MCMC samples for the latent variable home using the proposed distribution.
4. (Bonus, 20 points) Set the parameters as $\sigma = 0.005, 0.05, 0.5$ and $t = 1, 5, 20, 50$, and:
 - Plot the trace plot of the burning phase and the MCMC samples for the latent variable home using proposal distributions with different σ and t .
 - Estimate the rejection ratio for each parameter setting, report your results in a table.
 - Comment on the results. Which parameter setting worked the best for the algorithm?
 - Use the results from the optimal parameter setting
 - plot the posterior histogram of home from the MCMC samples
 - plot the estimated attacking strength $\mathbb{E}_{p(\theta, \eta|y)}[\text{att}_i]$ against the estimated defending strength $\mathbb{E}_{p(\theta, \eta|y)}[\text{def}_i]$ for each the team in one scatter plot. Please make sure to identify the team index of each point on your scatter plot.

You are NOT allowed to use any existing implementations of M-H in this problem. Please include all the required results (figures + tables) in your writeup PDF submission.

2 Variational Inference [40 points]

2.1 KL-Divergence

- (10 points) In this section we consider the Kullback-Leibler divergence

$$KL(p||q) = - \int p(x) \ln \frac{q(x)}{p(x)} dx$$

Evaluate $KL(p(x)||q(x))$ where $p(x)$ and $q(x)$ are:


- Two scalar Gaussians $p(x) = \mathcal{N}(x; \mu, \sigma^2)$ and $q(x) = \mathcal{N}(x; m, s^2)$
 - two multivariate Gaussians $p(\mathbf{x}) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}, \boldsymbol{\Sigma})$ and $q(\mathbf{x}) = \mathcal{N}(\mathbf{x}; \mathbf{m}, \mathbf{s})$
- (30 points) Consider two random variables x and y which are related as:

$$t = x + v, v \sim \mathcal{N}(0, 1)$$

$$y = \begin{cases} 1 & \text{if } t > 0, \\ 0 & \text{otherwise} \end{cases}$$

We also have prior information that $x \sim \mathcal{N}(0, 1)$ and we receive one measurement $y = 1$.


Using the factor of the model here:



where

$$\begin{aligned} f_x(x) &= \mathcal{N}(x; 0, 1) \\ f_{xt}(x, t) &= \mathcal{N}(t; x, 1) \\ f_{ty}(t, y) &= \delta(y - \text{sign}(t)) \end{aligned}$$

(Another equivalent factor graph is



$$\begin{aligned} f_x(x) &= \mathcal{N}(x; 0, 1) \\ f_v(v) &= \mathcal{N}(v; 0, 1) \\ f_{vxt}(v, x, t) &= \delta(t - v - x) \\ f_{ty}(t, y) &= \delta(y - \text{sign}(t)) \end{aligned}$$

Since we are not interested in v itself (we do not want to infer $p(v|y)$, for instance), we opt for the other formulation, since it is more compact and therefore requires fewer messages in the message passing. Both formulations will result in exactly the same expression for $p(x|y)$.)

- (a) Use message passing with moment matching with pen and paper to compute $p(x|y = 1)$.

Hint: The mean and variance for the half-normal distribution

$$\mathcal{HN}(x; \sigma^2) = 2\mathcal{N}(x; 0, \sigma^2), x > 0$$

are

$$\mathbb{E}[x] = \frac{\sigma\sqrt{2}}{\sqrt{\pi}}, \text{Var}[x] = \sigma^2 \left(1 - \frac{2}{\pi}\right)$$

- (b) Write a program to verify the calculated distribution using importance sampling.

3 Consistency of Lasso [20 points]

Before going into graphical lasso, let's first consider a linear regression problem, with covariates $X \in \mathbb{R}^{n \times p}$ and response $y \in \mathbb{R}^n$. In the high-dimensional setting $n \ll p$, the ordinary least squares (OLS) regression will not generalize, so we need a regularized least squares as our model. We consider one of the most prominent regularized regression models, namely the Lasso, as our main tool in this homework problem. Lasso estimates the regression coefficients as

$$\hat{\beta}^{\text{lasso}} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|y - X\beta\|_2^2 + \lambda \|\beta\|_1 \quad (1)$$

where λ is a hyperparameter that governs the strength of the regularization and controls the sparsity of the coefficients identified.

The attachment contains the training data $(X_{\text{train}}, y_{\text{train}})$ and the test data $(X_{\text{test}}, y_{\text{test}})$. You can use your favorite Lasso implementation, such as `sklearn.linear_model.Lasso` or `glmnet` in R. We will use mean squared error (MSE) as the main evaluation metric.

3.1 Warm-up (5 points)

Let's do some warm-ups.

- (a) First trial (2 points). Fit a Lasso model with $(X_{\text{train}}, y_{\text{train}})$ and test it with $(X_{\text{test}}, y_{\text{test}})$, report $\text{MSE}_{\text{train}}$ and MSE_{test} . You should observe a generalization gap.
- (b) Hyperparameter tuning (3 points). Tuning hyperparameters to improve the performance has seemingly become a controversial strategy nowadays.

Nonetheless, let's experiment with some choices of λ and check the performance. Please repeat the basic experiment above with 10 choices of λ s evenly spaced on a log scale from 10^{-5} to 10^5 . Report one plot showing both the $\text{MSE}_{\text{train}}$ and MSE_{test} as a function of λ .

3.2 Weak Irrepresentable Condition (5 points)

You should notice that there is always a generalization gap between training and testing, and the gap seems larger than what can be expected from the measurement errors. Is it some property of the data that has trapped us from closing the generalization gap? The answer is yes.

The data is indeed generated from a linear Gaussian model as follows:

$$y^{(i)} = X^{(i)}\beta^* + \epsilon^{(i)}, \epsilon^{(i)} \sim \mathcal{N}(0, 1), i = 1, \dots, n \quad (2)$$

However, with some caveats:

- Only q covariates ($q < p$) are active, i.e., associated with the response. In other words, the true $\beta^* \in \mathbb{R}^p$ has q nonzeros.
- For each active covariate j , $\beta_j^* \sim \mathcal{U}(0, 5)$ and $X_j^{(i)} \sim \mathcal{N}(0, 1)$ for $i = 1, \dots, n$.
- What about the rest of the $p - q$ features? In X_{train} , they are duplicates of the active covariates. However, X_{test} is not constructed as so.

Let X_a be the active covariates of X_{train} and X_b be the remaining. We now offer a theoretical tool: if Lasso can correctly identify the active covariates, then

$$|C_{ba}C_{aa}^{-1}\mathbf{1}| < 1 \quad (3)$$

where $C_{ba} = \frac{1}{n}X_b^T X_a$, $C_{aa} = \frac{1}{n}X_a^T X_a$, $\mathbf{1}$ denotes a vector of ones, and the inequality holds element-wise.

Show that Lasso cannot correctly identify the active covariates with the data generated as above. It is not required, but please refer to (?) for further information.

3.3 Improving the Performance (10 points)

It looks like a vanilla Lasso will never solve our problem. Fortunately, we have more knowledge of data. In this section, we will design better methods that take advantage of the knowledge of the data and hopefully get better MSE.

For all the following two questions, please emphasize the design rationale of the method. Regarding empirical performance, please report it in a single plot showing both $\text{MSE}_{\text{train}}$ and MSE_{test} as a function of the hyperparameter. You do not have to stick with Lasso, but please limit yourself within, vaguely, the family of regularized least squares. The grading will significantly value the rationale of the methods than the actual empirical performance since random trial-and-error may also lead to good performance due to the simplicity of the data.

- (a) Heterogeneity of the samples (5 points). There are rarely truly iid data in the real world and the heterogeneity of the samples often create some spurious signals identified by the model. We offer an extra piece of knowledge of the data:

- For the remaining $p - q$ covariates of X_{train} , when we create them by duplicating the active covariates, we did not duplicate for every sample, but only 90% of the samples.

Please take advantage of this message, design a method, test it, and report the performance. Please be creative, but if one needs some inspiration, (?) may offer some.

- (b) Structure of the features (5 points). Another perspective is to take advantage of the knowledge of features, which can often be introduced by some intuitive understanding of the problem in reality. We offer another piece of knowledge of the data:

- If the i th covariate is active and the j th covariate is its duplicate, then $i < j$.

Please take advantage of this message, design a method, test it, and report the performance. Please be creative, but if one needs some inspiration, "stepwise selection" may offer some.

References

- [1] 2013–14 Premier League. URL https://en.wikipedia.org/w/index.php?title=2013%E2%80%9314_Premier_League&oldid=1239520943. Page Version ID: 1239520943.
- [2] G. Baio and M. Blangiardo. Bayesian hierarchical model for the prediction of football results. 37(2):253–264. ISSN 0266-4763, 1360-0532. doi: 10.1080/02664760802684177. URL <https://www.tandfonline.com/doi/full/10.1080/02664760802684177>.